

ALGOREP Project

Contents

1	Software architecture	1
1.1	Messenger	1
1.2	Receiver	2
1.2.1	MessageReceiver	2
1.2.2	ReceiverManager	3
1.2.3	ConsensusManager	4
1.2.4	ElectionManager	4
1.2.5	FailureManager	4
1.2.6	ClientManager	5
1.2.7	ReplManager	6
2	Tests	6
3	Performance measures	7
3.1	Election	7
3.2	Consensus	7

Gwenegán Bertho
Kévin Guillet
Mathieu Tammaro
Jose A. Henriquez Roa

November 1, 2021

1 Software architecture

This section shall go through all the design choices made throughout the development of the project. First, a description of Messenger class, which fully encapsulates the Open MPI API. Then, an overview of the MessageReceiver virtual class, followed by a description of the 5 classes inheriting from it, and the ReceiverManager class that takes care of managing all instantiations of the derived classes. And finally, a brief description of the Node and Client classes.

1.1 Messenger




Figure 1: Messenger class diagram

The messenger class, located in `src/include/messenger.hh`, is a complete encapsulation of the Open MPI API. All MPI related calls are through the member functions of this class. When it comes to making use of APIs, the design choice encapsulating them is fairly common and easily justified. And as with any design choice we have done this for the sake of decreasing the complexity of the overall architecture. Thus, we list some of the benefits of this encapsulation;

First, it allowed to safely learn to use the API by localizing any changes made to the program to a single class. Seen as no matter how much the manual of an API is read, it remains very likely during development that the use of a given functionality from the API that was previously used during the early stages might no longer be desired during later stages. For instance during the development of this project a similar scenario was encountered where an ambiguity in the MPI standard led to an incorrect assumption on the behavior of a function; In the 4'th version of the Message-Passing Interface Standard in section 11.9.3 detailing the client routines to use when implementing a client/server model with MPI the standard states that:

*“If the port exists, but does not have a pending MPI_COMM_ACCEPT , the connection attempt will eventually time out after an implementation-defined time, **or** succeed when the server calls MPI_COMM_ACCEPT .In the case of a time out, MPI_COMM_CONNECT raises an error of class MPIERR_PORT .”*

This statement differs in the Open MPI documentation, where the description of the function MPI.Comm.Connect states that:

“The MPI_Comm_connect call must only be called after the MPI_Comm_accept call has been made by the MPI job acting as the server.”.

As was found during the actual use of the MPI_Comm_Connect the the highlighted 'or' in the quote from the standard turned out to be an exclusive 'or' where the choice to either implement the timeout functionality or allow the 'connect' call to complete upon 'accept' being called on the server was left to the implementation. This meant that there would be an intrinsic need to somehow handle race conditions from the client, due to the fact that a failed 'connect' call would eventually timeout the client, which would then force the client to exit with an MPI_ERR_PORT error. This was eventually resolved by making the clients synchronize themselves though the use of an external file, located in `etc/turn.txt`. In this scenario this incorrect assumption led to the unforeseen implementation of a client synchronization mechanism. However, if for instance an incorrect assumption had been made on the use of the MPI_Recv, causing the need for this call to be switched to the non-blocking variant MPI_Irecv. Then having encapsulated the API would then limit the modification required to be made to a single class, instead of every line throughout the project where the function was used.

Another, use for encapsulating the API is to limit the dependency to said API. For instance, if there ever came a need to switch Open MPI to another MPI implementation then there would only be need to switch out the API calls on a single class instead of the entire project. And the final use for this encapsulation is to simply hide as much logic related to the API as possible from outside classes. And so be able to work on other sections of the project without having to keep in mind the usage of the API and its functionalities. And so, reduce the overall complexity of the design.

1.2 Receiver

The following sections are all of the classes related to the `MessageReceiver` class, including this one.

1.2.1 MessageReceiver

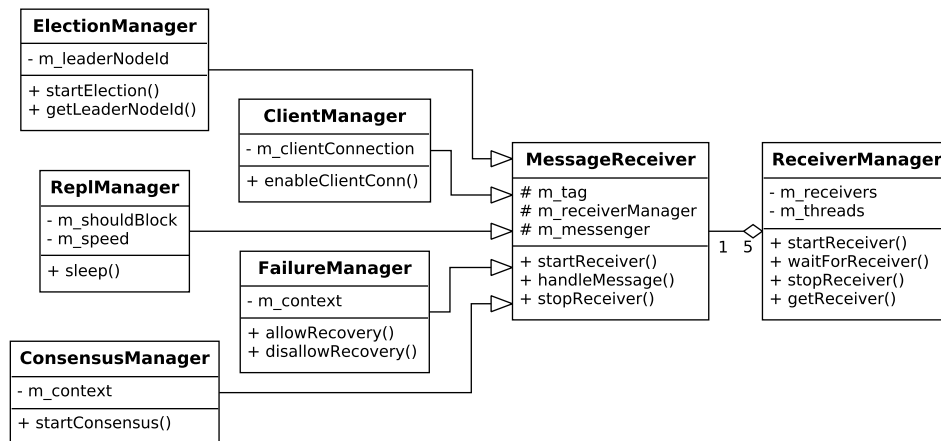


Figure 2: MessageReceiver class diagram

The `MessageReceiver` class is the virtual class at the center of the class architecture of the whole project. This class encapsulates all logic related to receiving and handling all messages from a single specified tag. This tag is specified in the constructor and is stored in the protected `m_tag` member variable. For the most part derived instances of this class, first specify the tag to handle, and simply start an infinite 'receive' loop when the `startReceiver()` function is called. In this loop the virtual `handleMessage()` member function is called after each new message. And so, derived classes only need to implement said `handleMessage()` function.

In this project there are five different message tags. And each of these tags possesses it's own respective set of *message codes*. All of which are defined in the `src/include/message-info.hh` header file. As mentioned, each of these tags are managed by their respective *Manager* inheriting from `MessageReceiver` class. The relations are as follows:

```

1  enum class MessageTag {
2      LEADER_ELECTION = 0,    // ElectionManager
3      CONSENSUS = 1,         // ConsensusManager
4      REPL = 2,              // ReplManager
5      FAILURE_DETECTION = 3,  // FailureManager
6      CLIENT = 4,            // ClientManager
7      SIZE = 5
8  };

```

Each of these five tag managers/receivers are instantiated once at the start of the program and are subscribed to the `ReceiverManager`. A description of this class and the five tag managers is given in the following sections.

1.2.2 ReceiverManager

As will be further detailed in the five subsequent sections each of the tag managers has some form of time sensitive logic to handle. In most of these sections there is need to wait for a given amount time. However, pausing the main thread for that given amount of time could possible entail an incorrect behavior on some other time sensitive section. And so for each of these time sensitive sections we were faced with one of two choices. Either manually implementing said time sensitive logic through the use of busy waiting functions or making use of multi-threading and pause current the thread to let the job scheduler switch out to an active thread. The choice was made to make use of multi-threading due to the number of time sensitive sections that were to be implemented which when coupled with the implementation of a specific busy waiting function for each of these sections would lead to a much greater amount of code to both write and maintain.

The `ReceiverManager` class is the one that manages these threads. One is created for each receiver and remains active until the program is issued to shutdown, handles to these threads are stored in the `m_threads` member variable. In addition to this the `ReceiverManager` also server as a container of each of the five receivers. Access from withing each of the receivers to other ones is done through the use of the `getReceiver()` member function of this class.

1.2.3 ConsensusManager

The ConsensusManager is a derived class from the MessageReceiver handling all messages tagged with `MessageTag::CONSENSUS`. This class only exposes a single `startConsensus()` function. The implementation of this class is located in `src/manager/consensus-manager.cc`. As with all classes an effort was made to hide as much of the inner logic as possible and reduce coupling of this class to other ones.

- `startConsensus()`: This function starts an consensus round. For the consensus our implementation of the Multi-Paxos algorithm was used. This function implements the behavior of a Proposer in the Paxos algorithm. In other words, depending on whether the respective conditions are met, it broadcasts `PREPARE`, `PROMISE` and `ACCEPTED` (from the `ConsensusCode` enum) to the acceptors. While also waiting for the respective responses for the respective given amounts of time.
- `handleMessage()`: The ConsensusManager's implementation of this function handles the Acceptor (and Learner) side of the Paxos algorithm. Meaning that upon receiving the respective messages and checking that some additional conditions are met, it sends `PROMISE` and `ACCEPT` back to the proposer.

That said, our implementation of this algorithm is a straightforward one.

1.2.4 ElectionManager

The ElectionManager is the derived class from the MessageReceiver and handles all messages tagged with `MessageTag::LEADER_ELECTION`. When first starting the receive loop this manager will startup an election to set the initial leader. This class exposes the two `startElection()` and `getLeaderNodeId()` member functions.

- `startElection()`: As the name implies this function starts an election. For leader election the Bully algorithm was used. This function broadcast a `ELECTION` message to all other nodes and then waits for a victory message. If none is received for the allotted amount of time (and no response was heard back from a node with a higher id), then a `VICTORY` message will be broadcasted.
- `handleMessage()`: Upon receival of the two election codes list above the this function will either send back an `ALIVE` message, and possibly start an election (depending on the id of the source node). Or simply set the inner `m_leaderNodeId` member variable.
- `getLeaderNodeId()`: This function simply return the leader node id. This function is only used in the FailureManager for determining whether the current node should handle the recovery of a failed node.

1.2.5 FailureManager

The FailureManager is the derived class from the MessageReceiver that handles all messages tagged with `MessageTag::FAILURE_DETECTION`. In this case most of the logic run by this class is done without outer interaction from other classes in a `pingCheck()` thread created when the receive loop is first started. Aside from this one the class exposes the two `allowRecovery()` and

`disallowRecovery()`.

- `handleMessage()`: For the FailureManager this function handles four different codes presented in the following list, the last three being related to node recovery:
 1. **PING**: When a ping is received. Then the respective value in a 'last-seen' list of time points is updated for the source node. This list is then checked in the `pingCheck` thread.
 2. **STATE**: This failure code is used when sending the current state of the system to a recovering node. When received the node simply overwrites the log data with the presented one inside the message. Only the current leader is allowed to send this code.
 3. **STATE.UDPATED**: This is the answer the recovering node sends back to the source node after having finished copying the new log state.
 4. **RECOVRED**: And finally, this is the message code broadcasted to all nodes when the node having issued the **STATE** code receives a **STATE.UDPATED** from the recovering node. This and the two previous message codes are coupled with a 'recovery id' to ensure correct behavior even during a change of leader.
- `pingCheck()`: Failure detection is done through all-to-all heart-beating. This function is run on a separate thread it executes an infinite loop that both broadcasts pings and check the 'last-seen' list at regular intervals. When either a failure or recovery is detected an additional (detached) thread is spawned to handle either event. This is done in order to minimize any latency these events could generate on the `pingCheck` thread.
- `allowRecovery()`: This functions is called by the clientManager to notify to the failureManager that there is no current active client connection (which by extension means that the log will not change from this point on). And that any pending recovery can start from this point on.
- `disallowRecovery()`: Also called by the clientManager it will notify the failureManager about the opposite. It will also block until the current recovery is done or the allotted time runs out.

1.2.6 ClientManager

The ClientManager is the derived class from the MessageReceiver that handles all messages tagged with `MessageTag::CLIENT`. This class only exposes the `enableClientConn()` member function.

- `handleMessage()`: For the ClientManager this function handles two codes **REPLICATE** and **DISCONNECT**. These respectively start a consensus round or notify the manager to disconnect the client. Once the consensus round has finished the manager responds back with **SUCCESS** to the client.
- `enableClientConn()`: Only the leader is allowed to accept client connection this is due to the nature of Multi-Paxos. This function is called by the ElectionManager when the current node is chosen as the leader.

1.2.7 ReplManager

The ReplManager is the derived class from the MessageReceiver that handles all messages tagged with `MessageTag::REPL`. This manager is somewhat different than the others. A greater deal of effort was made minimize the coupling from this class to the rest of the project. This mainly been due to the fact that on it's own it implements a completely separate set of functionalities that are unrelated to all others. And so, the class architecture design was made to make the removal of this particular class as easy as possible if there ever came a need.

With this in mind, during development of this class we faced the choice to; either separate an repl process from the node cluster and implement an repl client that would read from standard input and forward the messages to the repl process on the server, or, make the repl manager read from commands from a file instead and keep the node cluster intact. Due to the aforementioned reason the later option was chosen. Finally, the class itself exposes a single `sleep()` member function.

- `handleMessage()`: This function handles the repl codes `START`, `SPEED_<LOW/MEDIUM/HIGH>`, `CRASH` and `RECOVER`.
- `sleep()`: Depending on the issued repl commands this function will either block the current thread, until `RECOVERY` is received. Or put the thread to sleep for a given amount of time depending on whether `SPEED_<LOW/MEDIUM/HIGH>` were received. This function is called by all receivers except this one right before their respective `handleMessage()` function.

2 Tests

Testing was exclusively done by hand through the use of debug messages rather than through a streamlined test suit. This is due to the multi-threaded nature of the application and to our belief that setting up such a system would require a too large of a portion of the imparted time. The following presents some of the debug messages that were used during development:

```
[0] [0] [1]: {"code":1}
[0] [0] [2]: {"code":1}
[0] [1] [2]: {"code":1}
[0] [2] [0]: {"code":2}
[0] [1] [0]: {"code":2}
[0] [1] [2]: {"code":1}
[0] [2] [1]: {"code":2}
[0] [2] [1]: {"code":2}
[0] [2] [0]: {"code":3}
[0] [2] [1]: {"code":3}
```

It shows an election between three nodes. The number between the first first set of brackets show the tag `MessageTag::LEADER_ELECTION` in this case and the following two sets of brackets respectively show the source node id and the destination node id of the message. What follows is the message code in json format. The mapping between the integer and what it means is as follows:

```

1  enum class LeaderElectionCode {
2      SHUTDOWN = 0,
3      ELECTION = 1,
4      ALIVE = 2,
5      VICTORY = 3
6  };

```

3 Performance measures

The following are performance statistics averaged accross 10 tests. The CPU in which these tests were performed was the AMD FX(tm)-8350 Eight-Core Processor. The single variable for these tests was the number of nodes used on the server side.




Figure 3: Election duration [left], Majority accept duration[right]

3.1 Election

The previous duration were retrieved by measuring the amount of time taken for a leader to be elected after having all nodes broadcast an **ELECTION** message. We believe the increase in the duration observed in between the systems made from seven to eight nodes is related to the fact that the tests were performed in an eight core CPU.

3.2 Consensus

The previous duration were retrieved by measuring the amount of time taken for the majority of acceptors to send back **ACCEPT** after the initial broadcast of the **PREPARE** message from the leader.